CLOUD COMPUTING CONCEPTS with Indranil Gupta (Indy)

LEADER

Lecture A

THE ELECTION PROBLEM

Why Election?

- Example 1: Your Bank account details are replicated at a few servers, but one of these servers is responsible for receiving all reads and writes, i.e., it is the leader among the replicas
 - What if there are two leaders per customer?
 - What if servers disagree about who the leader is?
 - What if the leader crashes?

 Each of the above scenarios leads to Inconsistency

More motivating examples

- Example 2: (A few lectures ago) In the sequencer-based algorithm for total ordering of multicasts, the "sequencer" = leader
- Example 3: Group of NTP servers: who is the root server?
- Other systems that need leader election: Apache Zookeeper, Google's Chubby
- Leader is useful for coordination among distributed servers

Leader Election Problem

- In a group of processes, elect a *Leader* to undertake special tasks
 - And *let everyone know* in the group about this Leader
- What happens when a leader fails (crashes)
 - Some process detects this (using a Failure Detector!)
 - Then what?
- Focus of this lecture: Election algorithm. Its goal:
 - 1. Elect one leader only among the non-faulty processes
 - 2. All non-faulty processes agree on who is the leader

System Model

- N processes.
- Each process has a unique id.
- Messages are eventually delivered.
- Failures may occur during the election protocol.

Calling for an Election

- Any process can call for an election.
- A process can call for at most one election at a time.
- Multiple processes are allowed to call an election simultaneously.
 - All of them together must yield only a single leader
- The result of an election should not depend on which process calls for it.

Election Problem, Formally

- A run of the election algorithm must always guarantee at the end:
 - Safety: For all non-faulty processes p: (p's elected = (q: a particular non-faulty process with the best attribute value) or Null)
 - Liveness: For all election runs: (election run terminates)
 & for all non-faulty processes p: p's elected is not Null
- At the end of the election protocol, the non-faulty process with the <u>best (highest)</u> election attribute value is elected.
 - Common attribute: leader has highest id
 - Other attribute examples: leader has highest IP address, or fastest cpu, or most disk space, or most number of files, etc.

Next

• A classical leader election algorithm.