# CLOUD COMPUTING CONCEPTS with Indranil Gupta (Indy)

## MUTUAL EXCLUSION

Lecture B

DISTRIBUTED MUTUAL EXCLUSION

## System Model

- Before solving any problem, specify its System Model:
  - Each pair of processes is connected by reliable channels (such as TCP).
  - Messages are eventually delivered to recipient, and in FIFO (First In First Out) order.
  - Processes do not fail.
    - Fault-tolerant variants exist in literature.

#### Central Solution

- Elect a central master (or leader)
  - Use one of our election algorithms!
- Master keeps
  - A queue of waiting requests from processes who wish to access the CS
  - A special **token** which allows its holder to access CS
- Actions of any process in group:
  - enter()
    - Send a request to master
    - Wait for token from master
  - exit()
    - Send back token to master

#### Central Solution

- Master Actions:
  - On receiving a request from process Pi
    if (master has token)
    Send token to Pi
    else

Add Pi to queue

• On receiving a token from process Pi

```
if (queue is not empty)
```

Dequeue head of queue (say Pj), send that process the token

else

Retain token

## Analysis of Central Algorithm

- Safety at most one process in CS
  - Exactly one token
- Liveness every request for CS granted eventually
  - With *N* processes in system, queue has at most *N* processes
  - If each process exits CS eventually and no failures, liveness guaranteed
- FIFO Ordering is guaranteed, in order of requests received at master

## Analyzing Performance

Efficient mutual exclusion algorithms use fewer messages, and make processes wait for shorter durations to access resources. Three metrics:

- **Bandwidth**: the total number of messages sent in each enter and exit operation.
- *Client delay*: the delay incurred by a process at each enter and exit operation (when *no* other process is in, or waiting)

(We will prefer mostly the enter operation.)

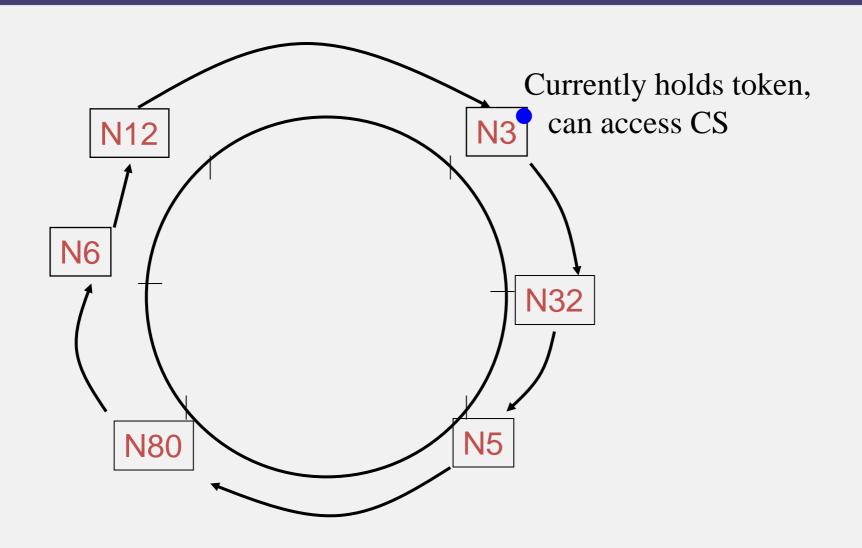
• Synchronization delay: the time interval between one process exiting the critical section and the next process entering it (when there is *only one* process waiting)

## Analysis of Central Algorithm

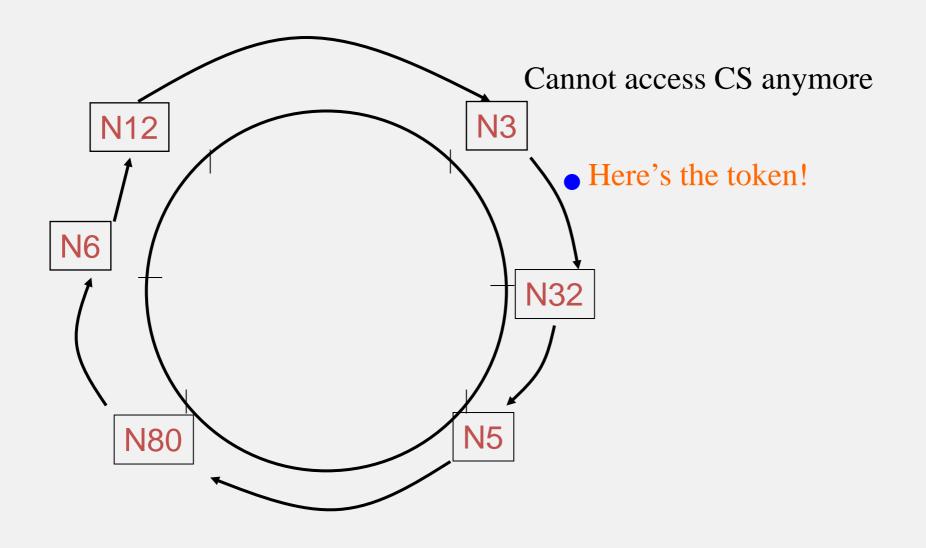
- **Bandwidth**: the total number of messages sent in each *enter* and *exit* operation.
  - 2 messages for enter
  - 1 message for exit
- *Client delay*: the delay incurred by a process at each enter and exit operation (when *no* other process is in, or waiting)
  - 2 message latencies (request + grant)
- Synchronization delay: the time interval between one process exiting the critical section and the next process entering it (when there is *only one* process waiting)
  - 2 message latencies (release + grant)

#### But...

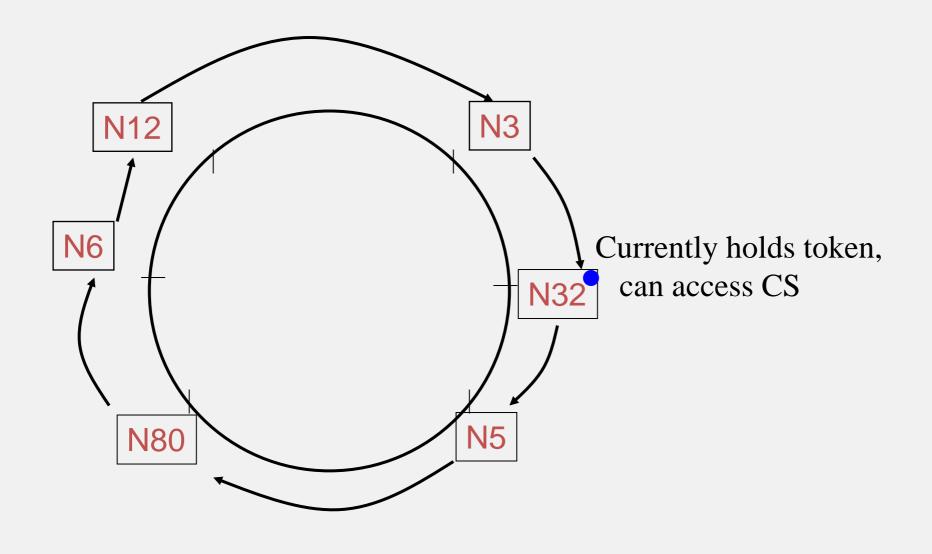
• The master is the performance bottleneck and SPoF (single point of failure)



Token: •



Token: •



Token: ●

- N Processes organized in a virtual ring
- Each process can send message to its successor in ring
- Exactly 1 token
- enter()
  - Wait until you get token
- exit() // already have token
  - Pass on token to ring successor
- If receive token, and not currently in enter(), just pass on token to ring successor

## Analysis of Ring-based Mutual Exclusion

- Safety
  - Exactly one token
- Liveness
  - Token eventually loops around ring and reaches requesting process (no failures)
- Bandwidth
  - Per enter(), 1 message by requesting process but *N* messages throughout system
  - 1 message sent per exit()

## Analysis of Ring-Based Mutual Exclusion (2)

- Client delay: 0 to N message transmissions after entering enter()
  - Best case: already have token
  - Worst case: just sent token to neighbor
- Synchronization delay between one process' exit() from the CS and the next process' enter():
  - Between 1 and (N-1) message transmissions.
  - <u>Best case</u>: process in enter() is successor of process in exit()
  - Worst case: process in enter() is predecessor of process in exit()

#### Next

- Client/Synchronization delay to access CS still O(*N*) in Ring-Based approach.
- Can we lower this?