



# CLOUD COMPUTING CONCEPTS

---

with Indranil Gupta (Indy)

## STREAM PROCESSING

Lecture A

---

STREAM PROCESSING IN STORM

# WHAT WE'LL COVER

- Why Stream Processing
- Storm

# STREAM PROCESSING CHALLENGE

- Large amounts of data => Need for real-time views of data
  - Social network trends, e.g., Twitter real-time search
  - Website statistics, e.g., Google Analytics
  - Intrusion detection systems, e.g., in most datacenters
- Process large amounts of data
  - With latencies of few seconds
  - With high throughput

# MAPREDUCE?

- Batch Processing => Need to wait for entire computation on large dataset to complete
- Not intended for long-running stream-processing

# ENTER STORM

- Apache Project
- <https://storm.incubator.apache.org/>
- Highly active JVM project
- Multiple languages supported via API
  - Python, Ruby, etc.
- Used by over 30 companies including
  - Twitter: For personalization, search
  - Flipboard: For generating custom feeds
  - Weather Channel, WebMD, etc.

# STORM COMPONENTS

- Tuples
- Streams
- Spouts
- Bolts
- Topologies

# TUPLE

- An ordered list of elements
- E.g., < tweeter, tweet >
  - E.g., < “Miley Cyrus”, “Hey! Here’s my new song!” >
  - E.g., < “Justin Bieber”, “Hey! Here’s MY new song!” >
- E.g., < URL, clicker-IP, date, time >
  - E.g., < coursera.org, 101.201.301.401, 4/4/2014, 10:35:40 >
  - E.g., < coursera.org, 901.801.701.601, 4/4/2014, 10:35:42 >

Tuple

# STREAM


- Sequence of tuples
  - Potentially unbounded in number of tuples
- Social network example:
  - $\langle \text{"Miley Cyrus"}, \text{"Hey! Here's my new song!"} \rangle$ ,
  - $\langle \text{"Justin Bieber"}, \text{"Hey! Here's MY new song!"} \rangle$ ,
  - $\langle \text{"Rolling Stones"}, \text{"Hey! Here's my old song that's still a super-hit!"} \rangle$ , ...
- Website example:
  - $\langle \text{coursera.org}, 101.201.301.401, 4/4/2014, 10:35:40 \rangle$ ,
  - $\langle \text{coursera.org}, 901.801.701.601, 4/4/2014, 10:35:42 \rangle$ , ...





# SPOUT


- A Storm entity (process) that is a source of streams
- Often reads from a crawler or DB



# BOLT




- A Storm entity (process) that
  - Processes input streams
  - Outputs more streams for other bolts




# TOPOLOGY

- A directed graph of spouts and bolts (and output bolts)
- Corresponds to a Storm “application”



# TOPOLOGY

- Can have cycles if the application requires it



# BOLTS COME IN MANY FLAVORS

- Operations that can be performed
  - **Filter:** forward only tuples which satisfy a condition
  - **Joins:** When receiving two streams A and B, output all pairs (A,B) which satisfy a condition
  - **Apply/transform:** Modify each tuple according to a function
  - And many others
- But bolts need to process a lot of data
  - Need to make them fast

# PARALLELIZING BOLTS

- Have multiple processes (“tasks”) constitute a bolt
- Incoming streams split among the tasks
- Typically each incoming tuple goes to one task in the bolt
  - Decided by “**Grouping strategy**”
- Three types of grouping are popular

# GROUPING

- **Shuffle Grouping**
  - Streams are distributed evenly across the bolt's tasks
  - Round-robin fashion
- **Fields Grouping**
  - Group a stream by a subset of its fields
  - E.g., All tweets where twitter username starts with [A-M,a-m,0-4] goes to task 1, and all tweets starting with [N-Z,n-z,5-9] go to task 2
- **All Grouping**
  - All tasks of bolt receive all input tuples
  - Useful for joins

# STORM CLUSTER

- Master node
  - Runs a daemon called *Nimbus*
  - Responsible for
    - Distributing code around cluster
    - Assigning tasks to machines
    - Monitoring for failures of machines
- Worker node
  - Runs on a machine (server)
  - Runs a daemon called *Supervisor*
  - Listens for work assigned to its machines
- Zookeeper
  - Coordinates Nimbus and Supervisors communication
  - All state of Supervisor and Nimbus is kept here



# FAILURES

- A tuple is considered failed when its topology (graph) of resulting tuples fails to be fully processed within a specified timeout
- **Anchoring:** Anchor an output to one or more input tuples
  - Failure of one tuple causes one or more tuples to replayed

# API FOR FAULT-TOLERANCE (OUTPUTCOLLECTOR)

- **Emit**(tuple, output)
  - Emits an output tuple, perhaps anchored on an input tuple (first argument)
- **Ack**(tuple)
  - Acknowledge that you finish processing a tuple
- **Fail**(tuple)
  - Immediately fail the spout tuple at the root of tuple topology if there is an exception from the database, etc.
- Must remember to ack/fail each tuple
  - Each tuple consumes memory. Failure to do so results in memory leaks.

# SUMMARY

- Processing data in real-time a big requirement today
- Storm
  - And other sister systems, e.g., Spark Streaming
- Parallelism
- Application topologies
- Fault-tolerance