

Case Study: Finding Primes

UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

© 2018 L. V. Kale at the University of Illinois Urbana

Example: List of Primes

- Consider the problem of finding the first million prime numbers
 - We want to output an array containing the first million primes

Sequential Code:

```
x = 3; j = 1; primes[0] = 2;
while(j < 1000000){
    if(isPrime(x)){
        primes[j] = x;
        j++;
    }
    x = x + 2;
}
```

- Note that `parallel for` requires that the loop bounds need to be known before the loop starts
- So, how can we parallelize this?

Ideas for Parallelizing PrimesList

- We can have each thread explore the next unexplored odd integer beginning with 3
- Both x and j need to be protected because multiple threads want to read and write them
- We can use locks or atomic variables for this purpose

```
x = 3; j = 1; primes[0] = 2;
while(j < 1000000){
    if(isPrime(x)){
        primes[j] = x;
        j++;
    }
    x = x + 2;
}
```

Version 1: Parallelizing PrimesList

- We protect writes to j and x using atomic
- Note that between the time a thread starts testing for x and the time it increments x, x may have been changed by other threads
 - So, after finishing testing 3, a thread may start working on testing 21, if other threads have already taken the number in between

```
int x, j;  
x = 3; j = 1; primes[0] = 2;  
#pragma omp parallel  
while (j < 1000000) {  
    if (isPrime(x)) {  
        primes[j] = x;  
  
#pragma omp atomic  
        j++;  
    }  
  
#pragma omp atomic  
    x = x + 2;  
}
```

Does this work?

Version 1: Parallelizing PrimesList

- The problem is between the time we test the primality of a number (x), and the time that it executes the assignment statement, some other thread might have changed the value of x
- Also, two different threads may try to assign with the same value of j

```
int x, j;  
x = 3; j = 1; primes[0] = 2;  
#pragma omp parallel  
while(j < 10000000) {  
    if(isPrime(x)) {  
        primes[j] = x;  
  
#pragma omp atomic  
        j++;  
    }  
  
#pragma omp atomic  
    x = x + 2;  
}
```

Does this work?

Version 2: Parallelizing PrimesList

- A thread saves the current value of x and increases it by 2 in a single critical section
 - No other thread can interfere
- A thread atomically increments j but saves the new value in its private variable k

```
int x, j, myX, k;  
x = 3; j = 0; primes[0] = 2;  
#pragma omp parallel private(myX, k)  
while(j < 10000000) {  
    #pragma omp atomic capture  
    { myX = x; x = x + 2; }  
    if(isPrime(myX)) {  
        #pragma omp atomic capture  
        k = j++;  
        primes[k] = myX;  
    }  
}
```

Does this work?

Version 2: Parallelizing PrimesList

- This almost works
- What are the problems?
- Is the primes array sorted?
 - No, although it is almost sorted
 - Some threads might run ahead
- How do we stop after the first one million primes?
 - While one thread is working on testing the millionth prime, another thread might finish testing the next prime and add it to the list

```
int x, j, myX, k;
x = 3; j = 0; primes[0] = 2;
#pragma omp parallel private(myX, k)
while(j < 1000000) {
    #pragma omp atomic capture
    { myX = x; x = x + 2; }
    if(isPrime(myX)) {
        #pragma omp atomic capture
        k = j++;
        primes[k] = myX;
    }
}
```

Fixing this is left as an exercise for you

Parallelizing PrimesList: Ideas for Fixes

- Let the loop go further for a few more iterations
- How many?
 - Maybe $j < 1000000 + \text{numThreads}$?
- Sort the array at the end?
 - Too expensive
 - And it's mostly sorted

```
int x, j, myX, k;
x = 3; j = 0; primes[0] = 2;
#pragma omp parallel private(myX, k)
while(j < 1000000) {
    #pragma omp atomic capture
    { myX = x; x = x + 2; }
    if(isPrime(myX)) {
        #pragma omp atomic capture
        k = j++;
        primes[k] = myX;
    }
}
```