


# Example: Prefix Sum

Recursive Doubling with Barriers

UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

# Prefix Sum Problem

- Given array  $A[0..N-1]$ , produce  $B[N]$ , such that  $B[k]$  is the sum of all elements of  $A$  up to  $A[k]$



$B[3]$  is the sum of  $A[0]$ ,  $A[1]$ ,  $A[2]$ ,  $A[3]$   
But  $B[3]$  can also be calculated as  $B[2] + A[3]$


# Prefix Sum: A good Sequential Algorithm

- Data dependency from iteration to iteration
  - How can this be parallelized at all?

```
B[0] = A[0];  
for (i=1; i<N; i++)  
    B[i] = B[i-1] + A[i];
```

- It looks like the problem is inherently sequential, but theoreticians came up with a beautiful algorithm called recursive doubling or just parallel prefix

# Parallel Prefix: Recursive Doubling



N Data Items

P Processors

$N=P$

Log P Phases

P additions in each phase


P log P operations

Completes in  $O(\log P)$  time

# OpenMP Formulation for Parallel Prefix

- We don't have  $n$  processors
  - I.e., the number of threads will be much smaller than the size of the data array
- So, we will simulate  $n$  processors using  $p$  threads
- Notice that each phase of the computation must finish before the next phase of the computation starts
  - We will use OpenMP's **barrier** directive for that
- Basic description of actions in each phase with distance  $d$ 
  - Every "processor"  $i$  adds its value to the value held by a processor distance  $d$  away
  - Simulation:  $B[i+d] += B[i]$ , but you have to be careful to avoid dependencies
    - I.e., copy  $B[i]$  into a temporary variable at  $i+d$ , say  $C[i+d]$  and then add  $C[i]$  to  $B[i]$  for every  $i$
  - Note  $d$  doubles in every phase

# Parallel Prefix: Recursive Doubling



```
...
#pragma omp parallel for
    for (i=0; i<n; i++) {B[i]=A[i];}
    int d=1;
    while (d<n) // this loop will run for lg n steps
    {
        int i;
        #pragma omp parallel for
        for (i=d; i<n; i++) C[i]=B[i-d];

        #pragma omp parallel for
        for (i=d; i<n; i++) B[i]+=C[i];

        d*=2;
    }
...
```

Initialize B with values from A

C[k] temporarily stores the value that we want to add to B[k]


# Critique of Prefix Algorithm 1

- The sequential algorithm had  $n$  additions
- But the parallel algorithm is doing a total of  $n \cdot (\log n)$  additions
  - Although they are parallelized by  $p$  threads
  - This is an example of an algorithm that is not “work efficient”
- It uses  $\log n$  barriers, which are expensive operations
- Maybe a thread oriented approach will avoid the  $\log n$  factors



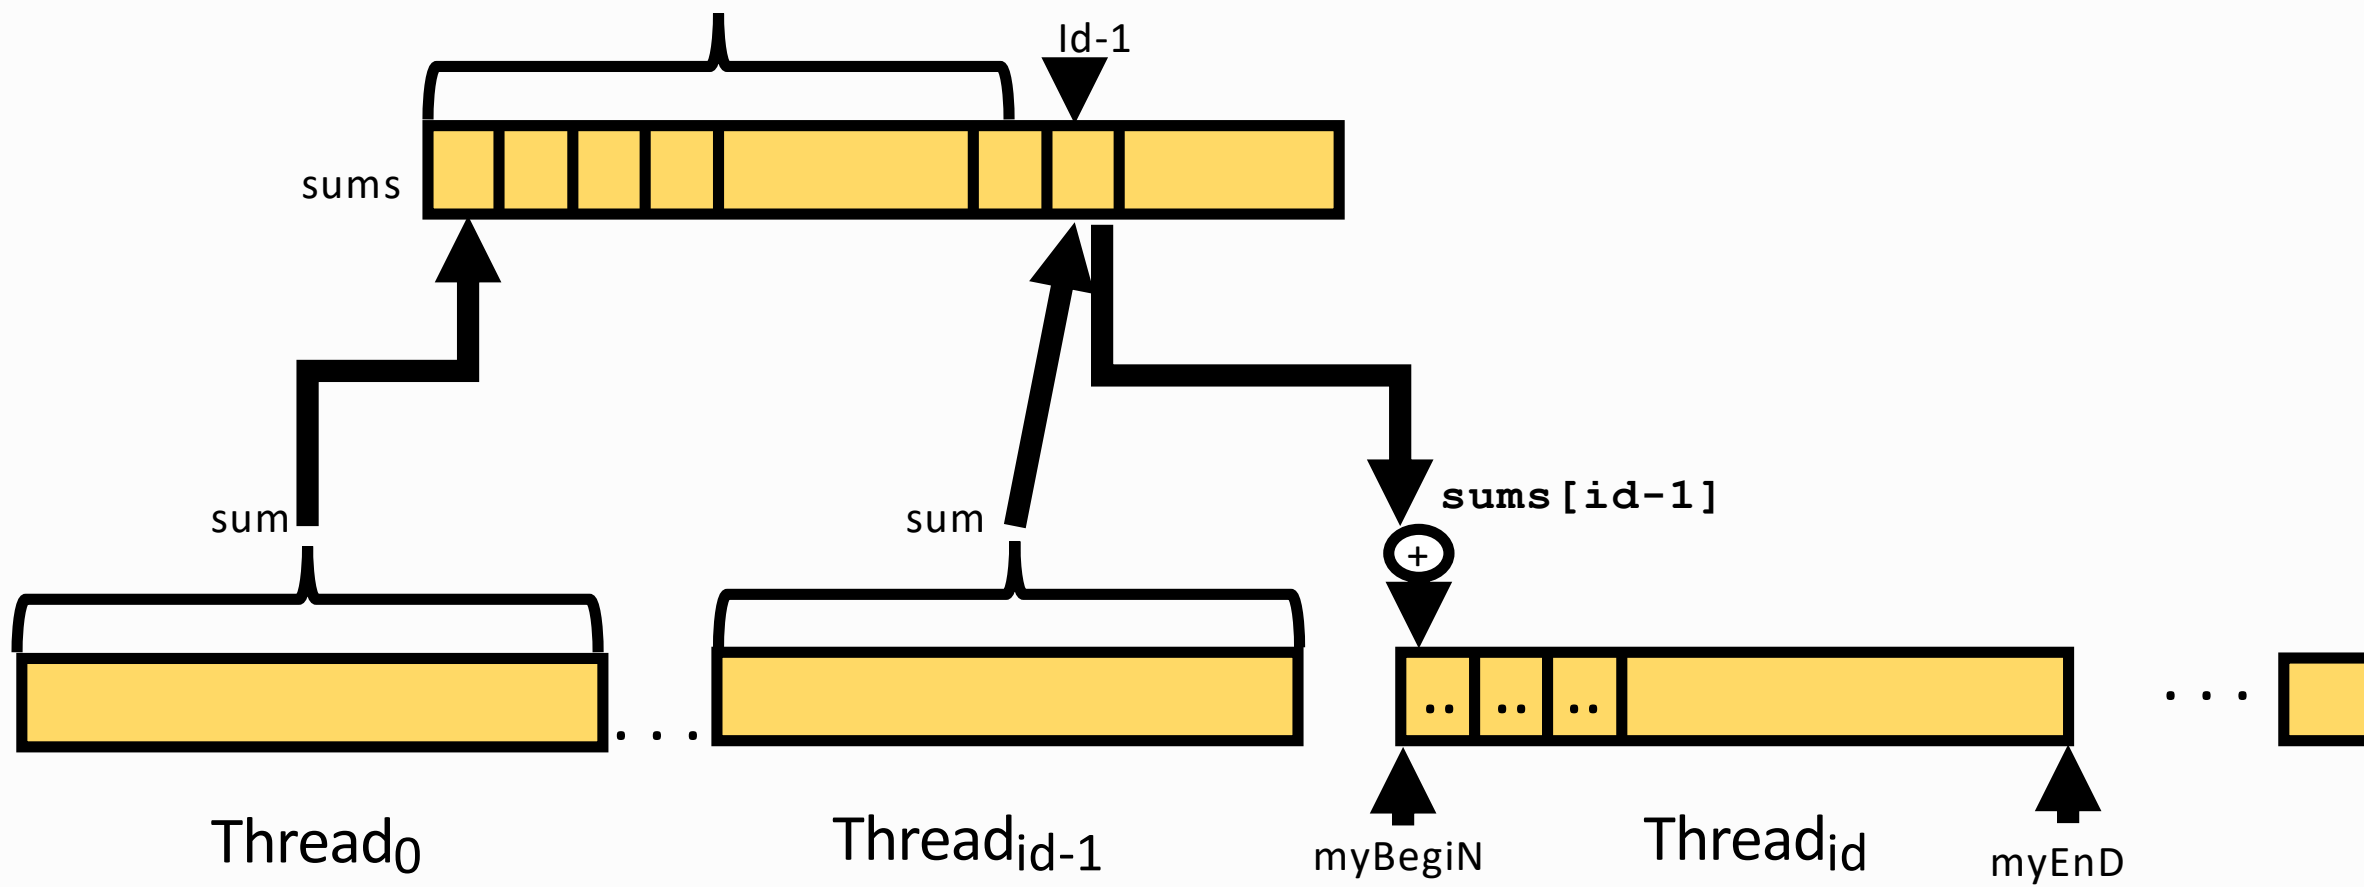
# Prefix Sum Algorithm 2: A Thread Oriented Approach

- What if we let each thread calculate prefix sum over its own range of array?
  - I.e., thread id is responsible for range  $B[\frac{n*id}{p} : \frac{n*(id+1)}{p} - 1]$
  - Id : my thread's serial number; p : total number of threads
  - Assuming n is a multiple of p
- But then each thread needs the sum of all numbers to its left



# Prefix Sum Algorithm 2: A Thread Oriented Approach

- What if we let each thread calculate prefix sum over its own range of array?
  - I.e., thread  $i$  is responsible for range  $B[\frac{n*id}{p} : \frac{n*(id+1)}{p} - 1]$   
\*id – my thread's serial number; \*p – total number of threads  
Assuming  $n$  is a multiple of  $p$
- But then each thread needs the sum of all numbers to its left
- If we are willing to double the amount of work, we can obtain this sum with a much smaller prefix sum problem of size  $p$ 
  1. First loop: every thread calculates sum  $s$  over its sub-range and copies  $s$  into a shared array called sums at **sums[id]**
  2. Calculate prefix sum of the sums array
    - **sums[id-1]** has the sum of all values to the left of thread numbered  $id$
  3. Second loop: every thread with serial number  $id$  calculates the prefix sum in array  $B$  using **sums[id-1]** and the values in  $A$



```
...
omp_set_num_threads(p);

#pragma omp parallel
{
    int id=omp_get_thread_num();
    int myBegin = (n*id)/p;
    int myEnd = min( (n*id+1)/p, n);
```

```
    int sum=0;
    for(int i=myBegin;i<myEnd;i++)
        sum+=B[i];
    sums[id]=sum;
```

```
    #pragma omp barrier
    #pragma omp single
    {
        for(int i=1;i<p;i++)
            sums[i]+=sums[i-1];
    }
```

```
    if(id>0) B[myBegin]+= sums[id-1]
    for(int i=myBegin+1;i<myEnd/p;i++)
        B[i]+=B[i-1];
}
```

Form Local sum

- Calculate Prefix sum of size p
- Sums [id] now contains the sum of values of all previous threads' ranges
- This can be done in parallel but it's not worth it

Complete the Prefix sum