


More Synchronization

Sequential Consistency, and the Flush Directive

UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

Memory Consistency Issues and OpenMP

- Let's start with an old “schematic” of a computer
- Instruction memory, CPU, and data memory
- Each arithmetic instruction is like: “add m1, m2, m3”
 - m1,m2,m3 are memory locations
 - Bring contents of m1 and m2 to ALU, add them up, and store the result in m3
- In addition, there are branch instructions

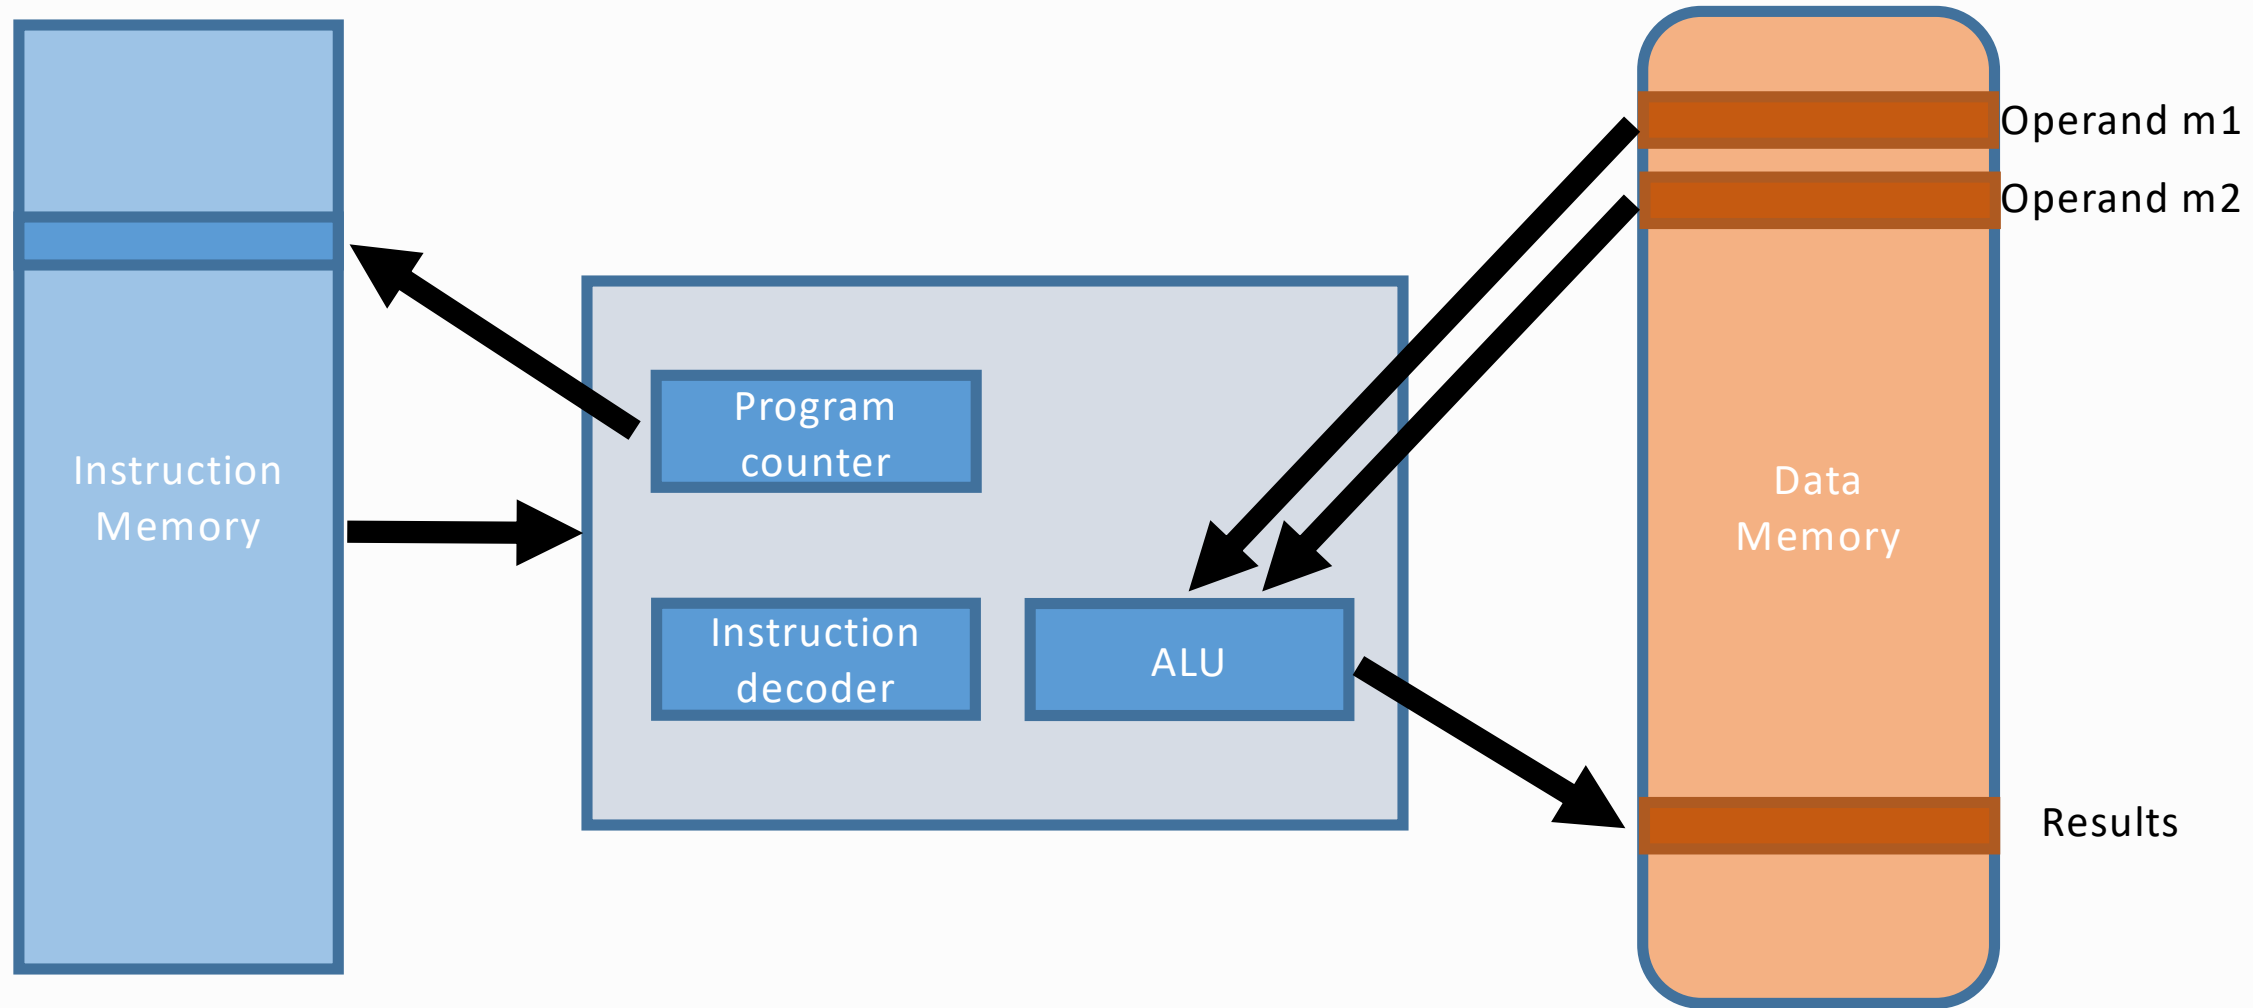


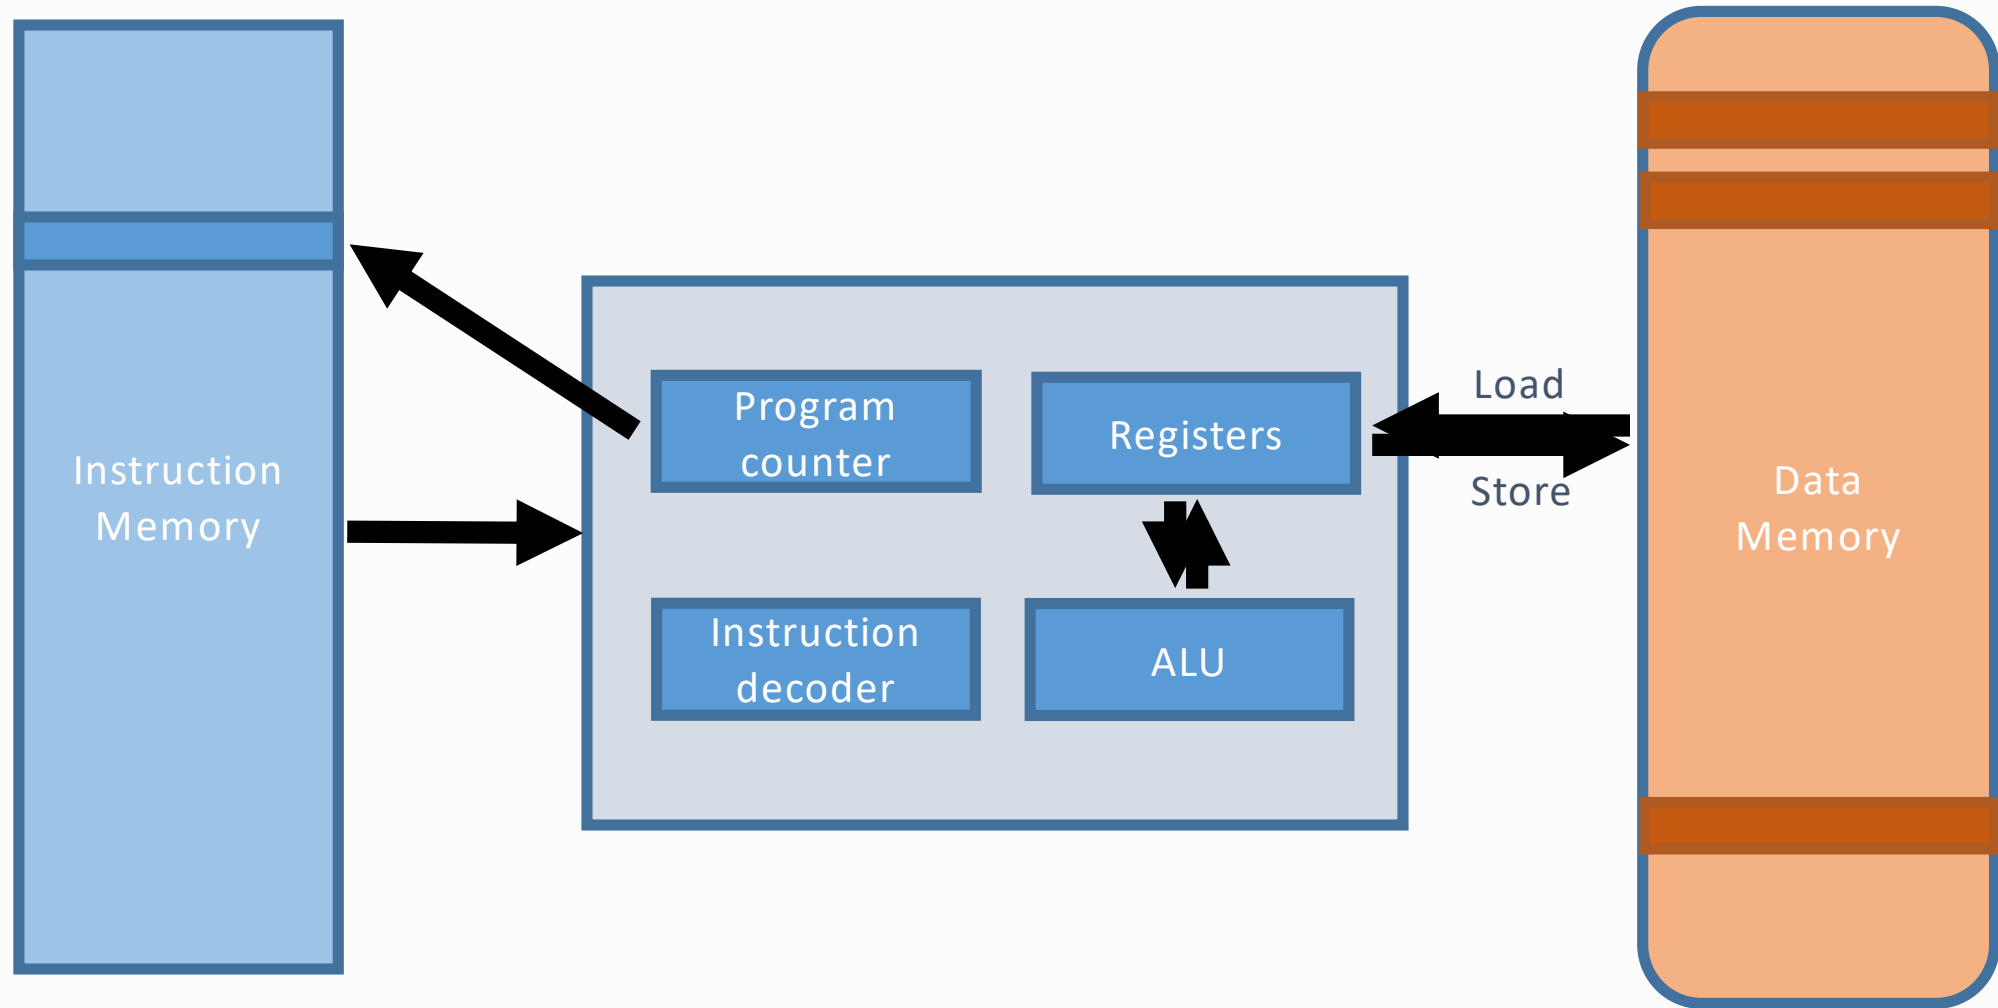
If Processors Remained That Simple ...

- If parallel computers with shared memory were built from such processors
 - Parallelism would be simpler
 - But they'd be slow
- Computer architects added caches, as we know
 - Some variables may exist in caches as well as memory
 - Cache coherence protocols, such as MESI for snoopy caches, handle data in caches
- But also, they added registers

Complexity of Real Computers: Registers

- Real computers became more complex before they became parallel
 - Remember: mostly to deal with slow memory
- Around 1985, the RISC revolution led to load-store architectures
 - Instructions either:
 - Do arithmetic/logical operations on registers, storing the result in registers,
 - Do a load from memory into register, OR
 - Do a store from register into memory





Complexity of Real Computers: Registers

- Real computers became more complex before they became parallel
 - Remember: mostly to deal with slow memory
- Around 1985, the RISC revolution led to load-store architectures
 - Instructions either:
 - Do arithmetic/logical operations on registers, storing the result in registers,
 - Do a load from memory into register, OR
 - Do a store from register into memory
- Variables can be stored in registers
 - But the association between variables and registers is loose
 - Not visible to the hardware
 - May be “spilled” lazily to memory via “write buffers”

Write Buffers

- For sequential processors, how long should a “Store” instruction take?
- You can speed it up by using a write buffer between the CPU and cache
- But this becomes problematic in parallel

Complexity: Compilers Can Reorder Statements

- Remember, compilers were written for sequential processors
- For example it may transform the first block of code below to the second block, by reordering statements
 - Why? Many reasons.. E.g. reducing the number of registers used, or eliminating a pipeline bubble

```
x = e1; // e1 doesn't contain y  
y = e2; // e2 doesn't contain x
```

```
y = e2; // e2 doesn't contain x  
x = e1; // e1 doesn't contain y
```


Complications with Parallel Processors

- Caches were handled using the extra hardware (snoopy cache controllers)
- But:
 - Data in registers
 - Data in “write buffers” (on its way to memory)
 - Instructions reordered by compiler
- When used with parallel processors, these tend to destroy our intuitive understanding of how parallel processors should behave, especially wrt memory
 - Notions of causality, program order, happens-before relation
- Our intuition is captured by the formalized notion of “sequential consistency”
 - Corresponds to the the simple picture of processor and memory we sketched earlier

Sequential Consistency

- This is a “desired property” of parallel programming systems
- The effect of executing a program consisting of k threads should be the same as some arbitrary interleaving of statements executed by each thread, executed sequentially

Modern processors do not satisfy sequential consistency!



Initially: x, Flag, are both 0

Thread 0:

x = 25;

Flag = 1;

Thread 1:

while (Flag == 0) ;

Print x;

What should get printed?

How to deal with lack of sequential consistency?

- Solution? Give up on registers and write buffers? No way!
- Various complicated processors specific synchronization primitives
- OpenMP provides simple machine-independent ***flush*** primitives

Producer Consumer Example

```
#pragma omp parallel {  
    if (omp_get_thread_num()==0) {  
        //Thread 0: Producer  
        Data = computeData();  
        flag = 1;  
  
    }  
    if(omp_get_thread_num()==1) {  
        //Thread 1: Consumer  
        while (flag==0) {  
            };  
        print Data;  
    }  
}
```


Producer Consumer Example

```
#pragma omp parallel {  
    if (omp_get_thread_num()==0) {  
        //Thread 0: Producer  
        Data = computeData();  
        #pragma omp flush (Data)  
        flag = 1;  
        #pragma omp flush (flag)  
    }  
    if(omp_get_thread_num()==1) {  
        //Thread 1: Consumer  
        while (flag==0) {  
            #pragma omp flush (flag)  
        };  
        print Data;  
    }  
}
```

Naming Variables in the Flush Directive

- If no variables are named, all are flushed
 - I.e., the processor waits until all writes from it in the past are visible and all registers are stored out to memory/cache (spilled, flushed)
- **flush** is especially useful for point-to-point synchronization
 - i.e. for one thread to signal to another (that some event has happened or some data is ready)