

# Gauss-Seidel

Dealing with Complicated Dependencies


UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

© 2018 L. V. Kale at the University of Illinois Urbana

# Gauss-Seidel Relaxation

Sequential pseudocode:

```
while (maxError > Threshold) {  
    Re-apply Boundary conditions  
    maxError = 0;  
    for i = 0 to N-1 {  
        for j = 0 to N-1 {  
            old = A[i, j]  
            A[i, j] = 0.2 * (A[i, j] + A[i, j-1] + A[i, j+1]  
                           + A[i+1, j] + A[i-1, j]) ;  
            if (|A[i, j]-old| > maxError)  
                maxError = |A[i, j]-old|  
        }  
    }  
}
```



For the same problem we solved using Jacobi Relaxation


No old-new arrays ...

Sequentially, how well does this work?

It works much better!

- Intuitively, the effect of boundary conditions spreads fast to other areas, compared with Gauss-Jacobi


How to parallelize this?



```
A[i, j] = 0.2 * (A[i,j] + A[i,j-1] + A[i,j+1]  
                + A[i+1,j] + A[i-1,j]) ;
```


# How Do We Parallelize Gauss-Seidel?

- Visualize the flow of values
- Not the control flow:
  - That goes row-by-row
- Flow of dependences: which values depend on which values?
- Does that give us a clue on how to parallelize?



# How Do We Parallelize Gauss-Seidel?


- Visualize the flow of values
- Not the control flow:
  - That goes row-by-row
- Flow of dependences: which values depend on which values?
- Does that give us a clue on how to parallelize?



```
for diagonal = 0 to 2*N-2 {  
  parallel loop over values in the diagonal  
  { i= .. ; j = ..;  
    old = A[i,j];  
    A[i, j] = ... ;  
    if (|A[i,j]-old| > maxError)  
      maxError = |A[i,j]-old|  
  }  
}
```

# Gauss-Seidel: parallelize each diagonal


- Performance is not so good. Why?
- Each thread is doing a different (shifting) section of rows.
  - Spatial locality and prefetch efficiency is affected
- Too fine grained a loop? There are  $2N$  parallel loops
- Other reasons? Implement and analyze with PAPI or perf tools



```
for diagonal = 0 to 2*N-2 {  
    parallel loop over values in the diagonal  
    { i= .. ; j = ..;  
      old = A[i,j];  
      A[i, j] = ... ;  
      if (|A[i,j]-old| > maxError)  
          maxError = |A[i,j]-old|  
    }  
}
```

# Parallelizing Gauss-Seidel

- Some ideas
  - Row decomposition, with pipelining




- Square over-decomposition
  - Assign many squares to a processor (essentially same?)

# Row decomposition with pipelining


# Columns =  $N/W$   
# Rows =  $P$

# Of Phases  
 $N/W + (P-1)$





# Row decomposition, with pipelining



# Red-Black Squares Method

- Red squares calculate values based on the black squares
  - Then black squares use values from red squares
  - Now red ones can be done in parallel, and then black ones can be done in parallel
- A “square” may be just a single point
  - Or it can be a  $k \times k$  tile of values
    - Each tile locally can do Gauss-Seidel computation
    - Faster convergence of Gauss-Seidel

