




CLOUD COMPUTING APPLICATIONS

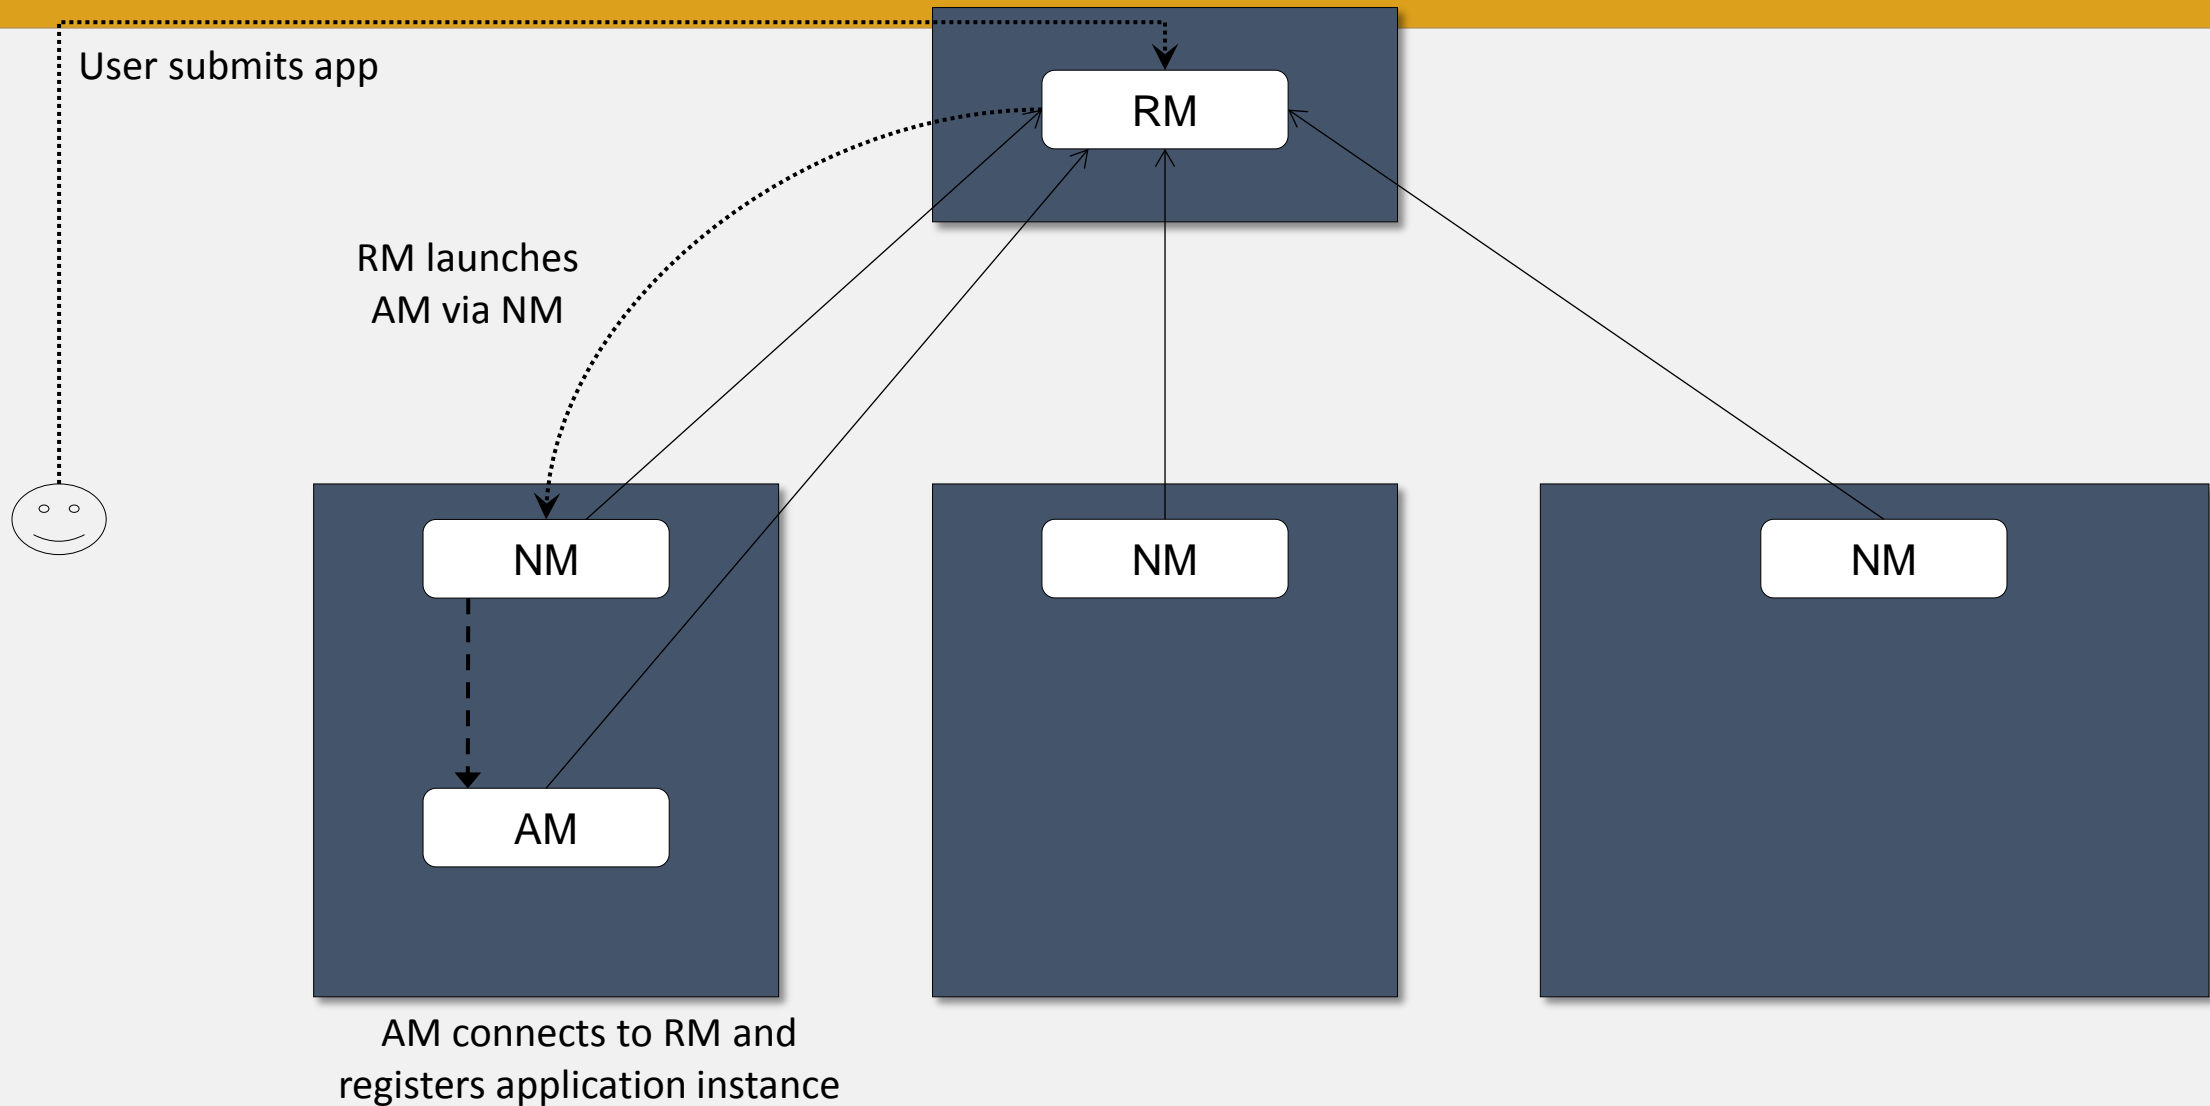
YARN: MapReduce on YARN in Diagram

Roy Campbell & Reza Farivar

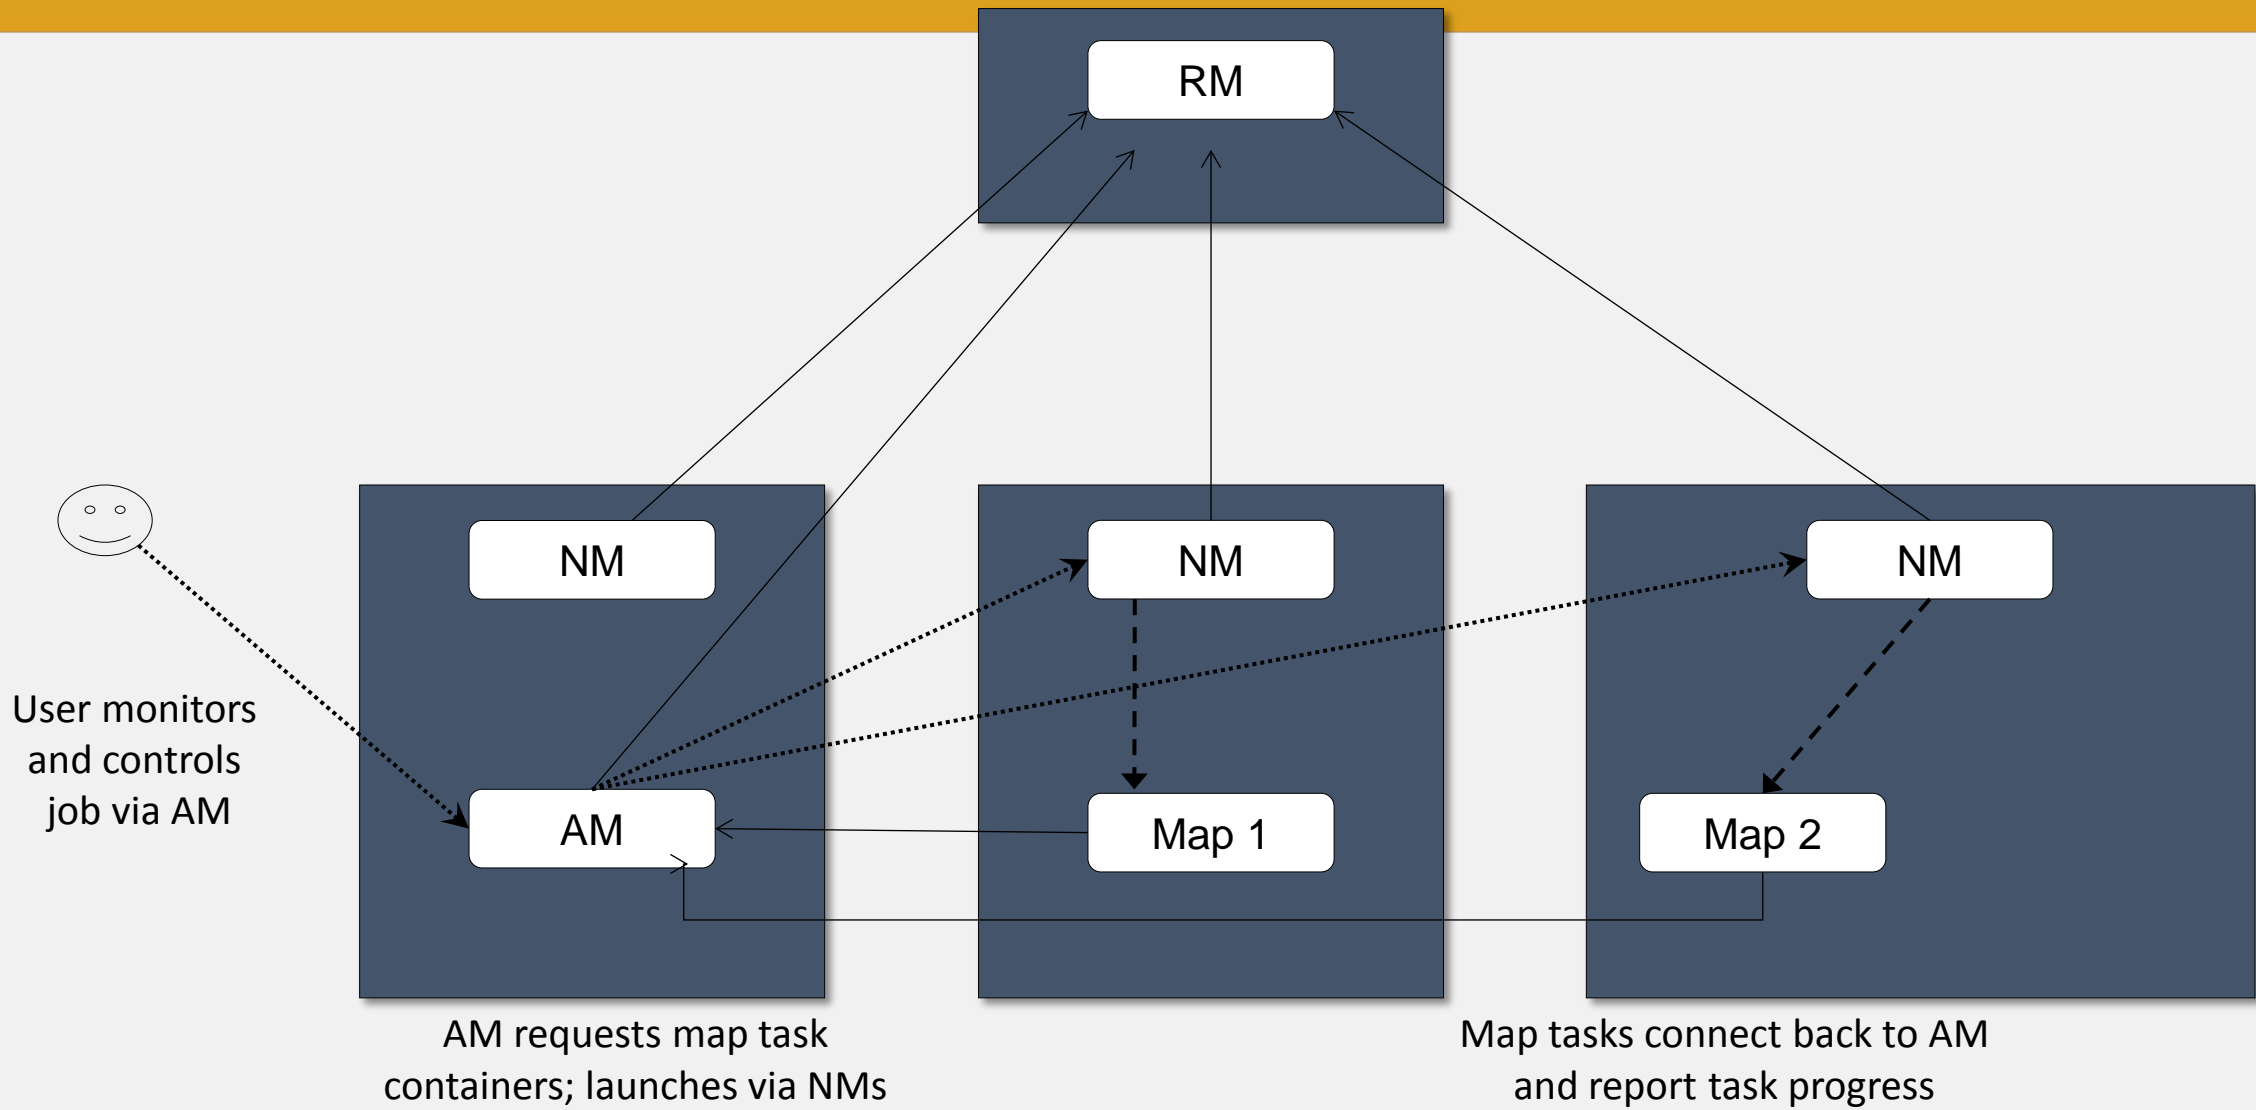
MapReduce on YARN



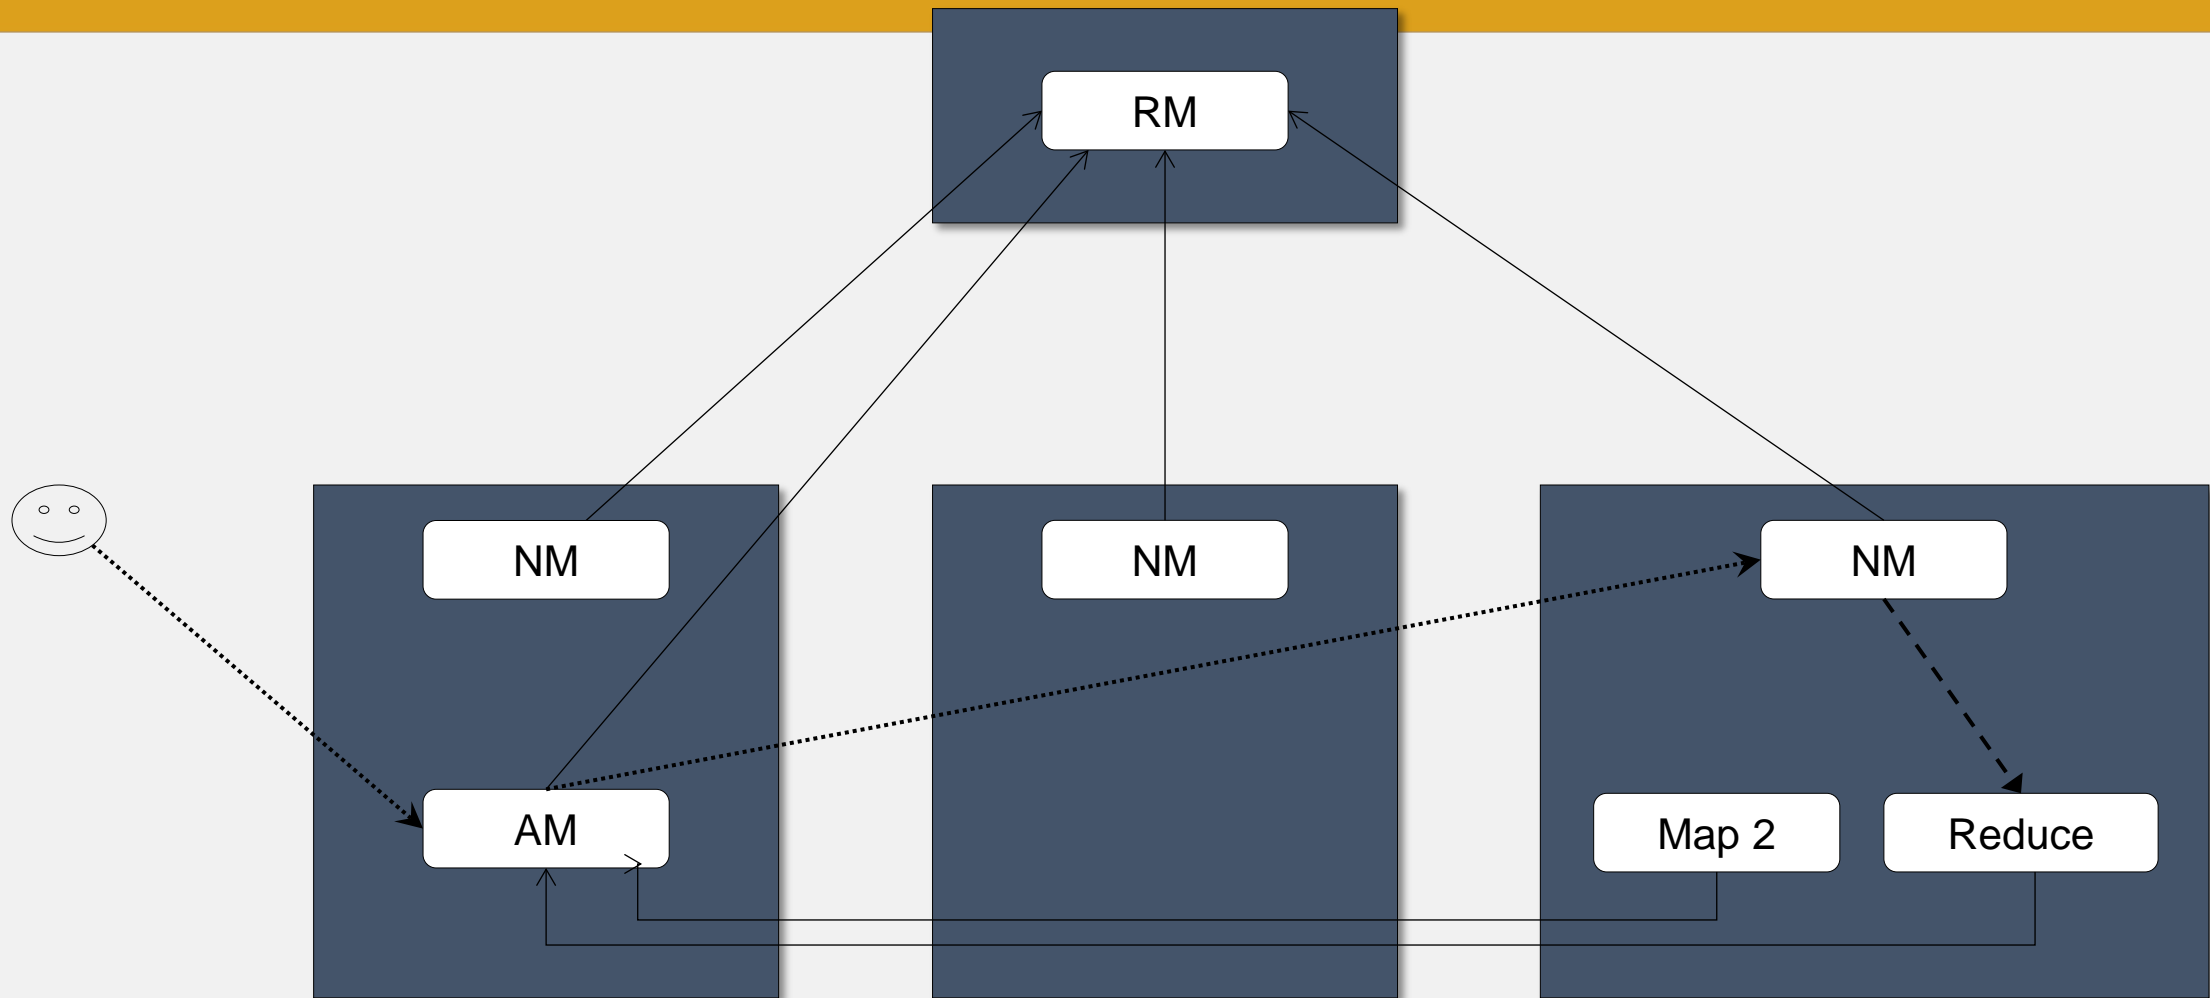
MapReduce on YARN



MapReduce on YARN

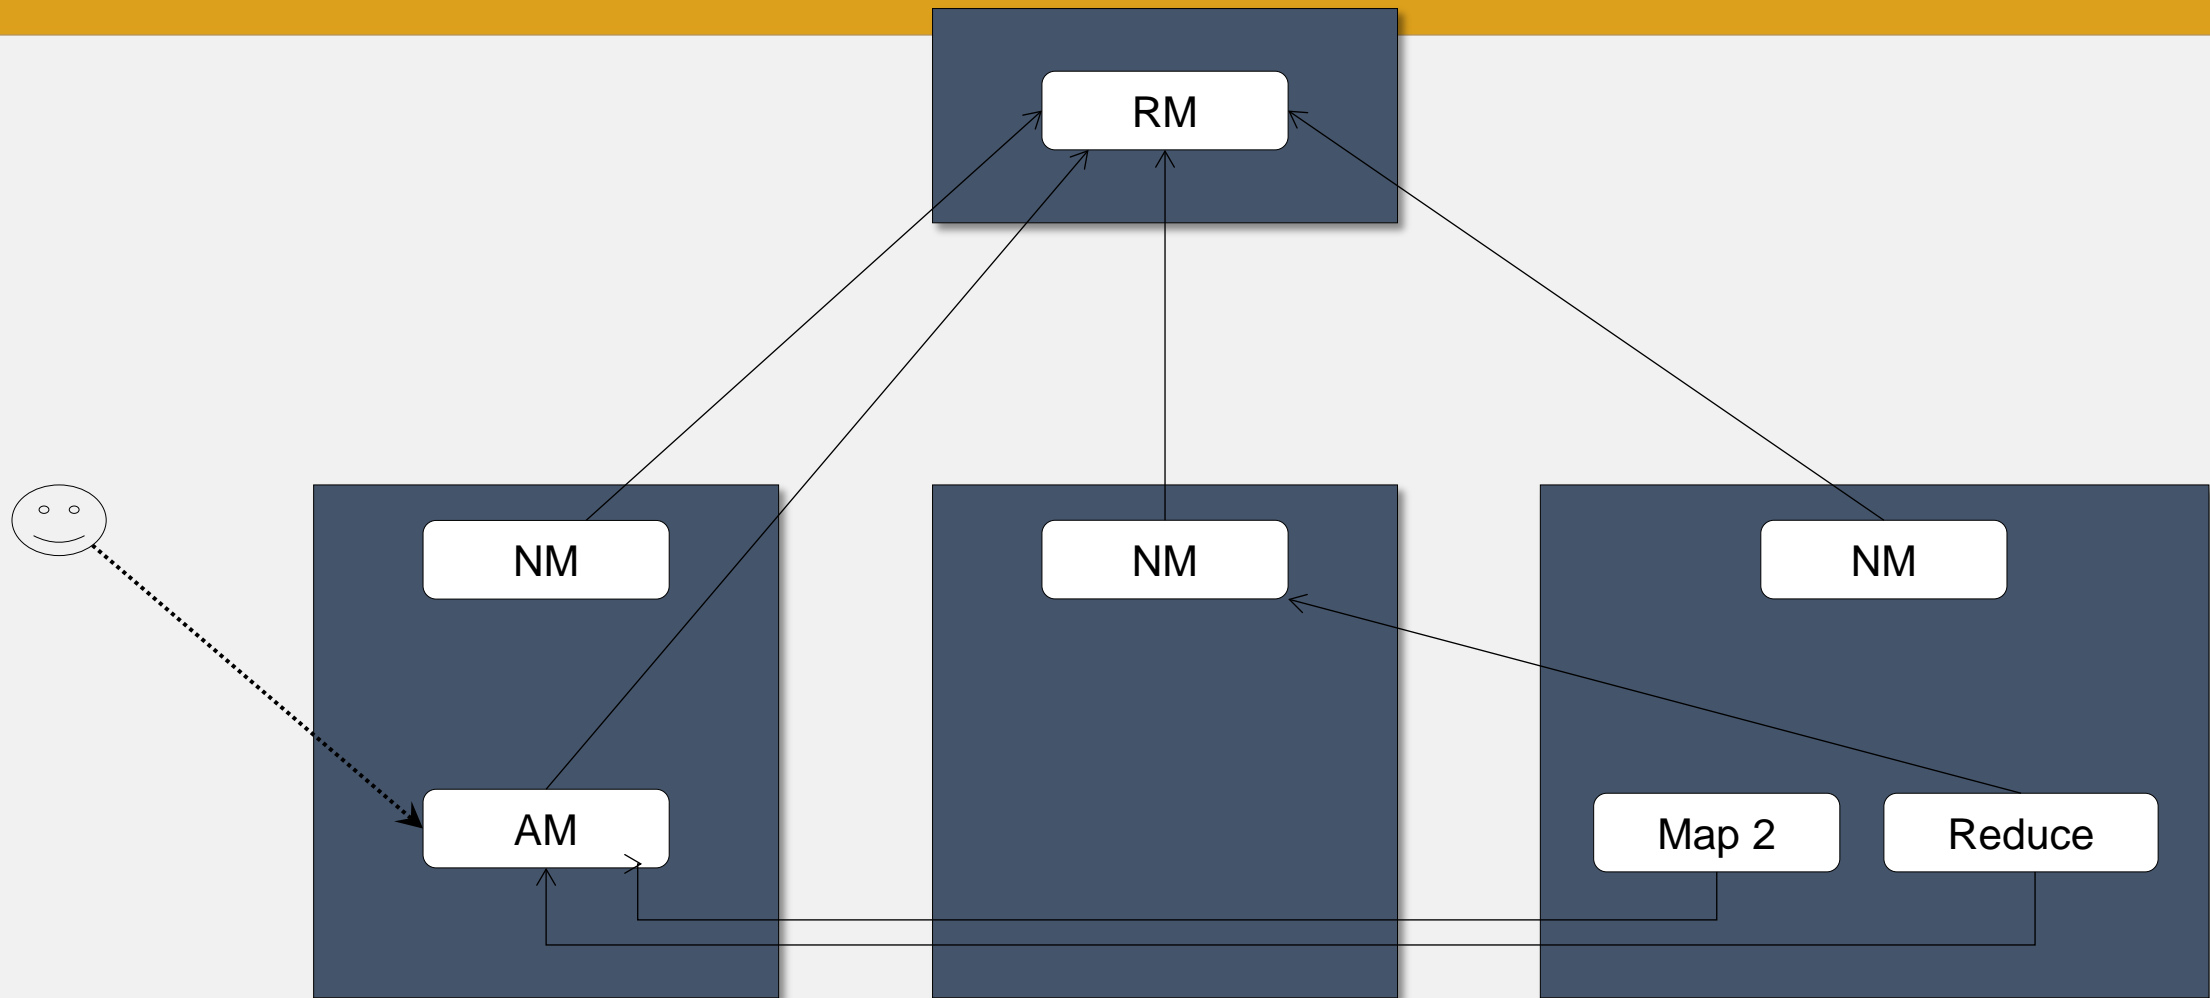


MapReduce on YARN



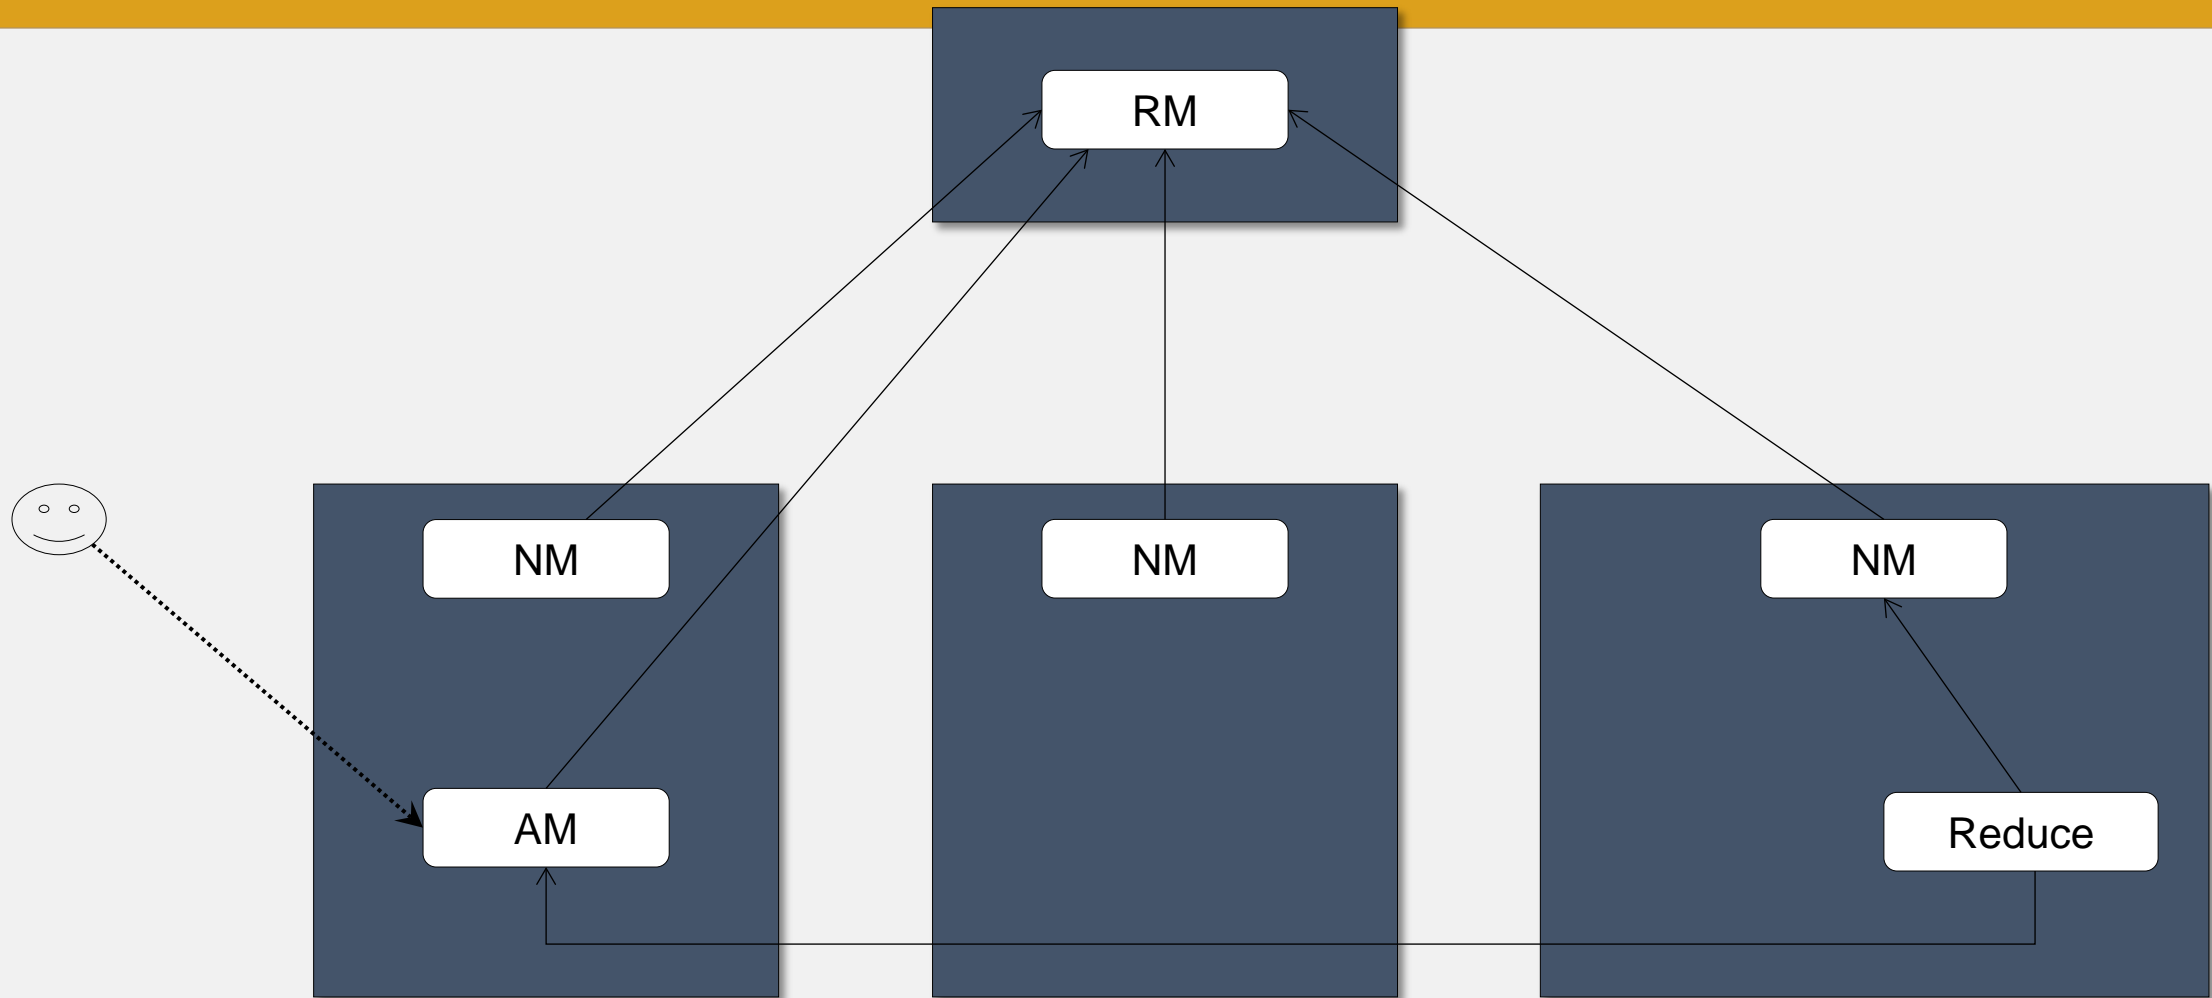
Map task 1 completes; AM requests container for reduce task and launches via NM

MapReduce on YARN



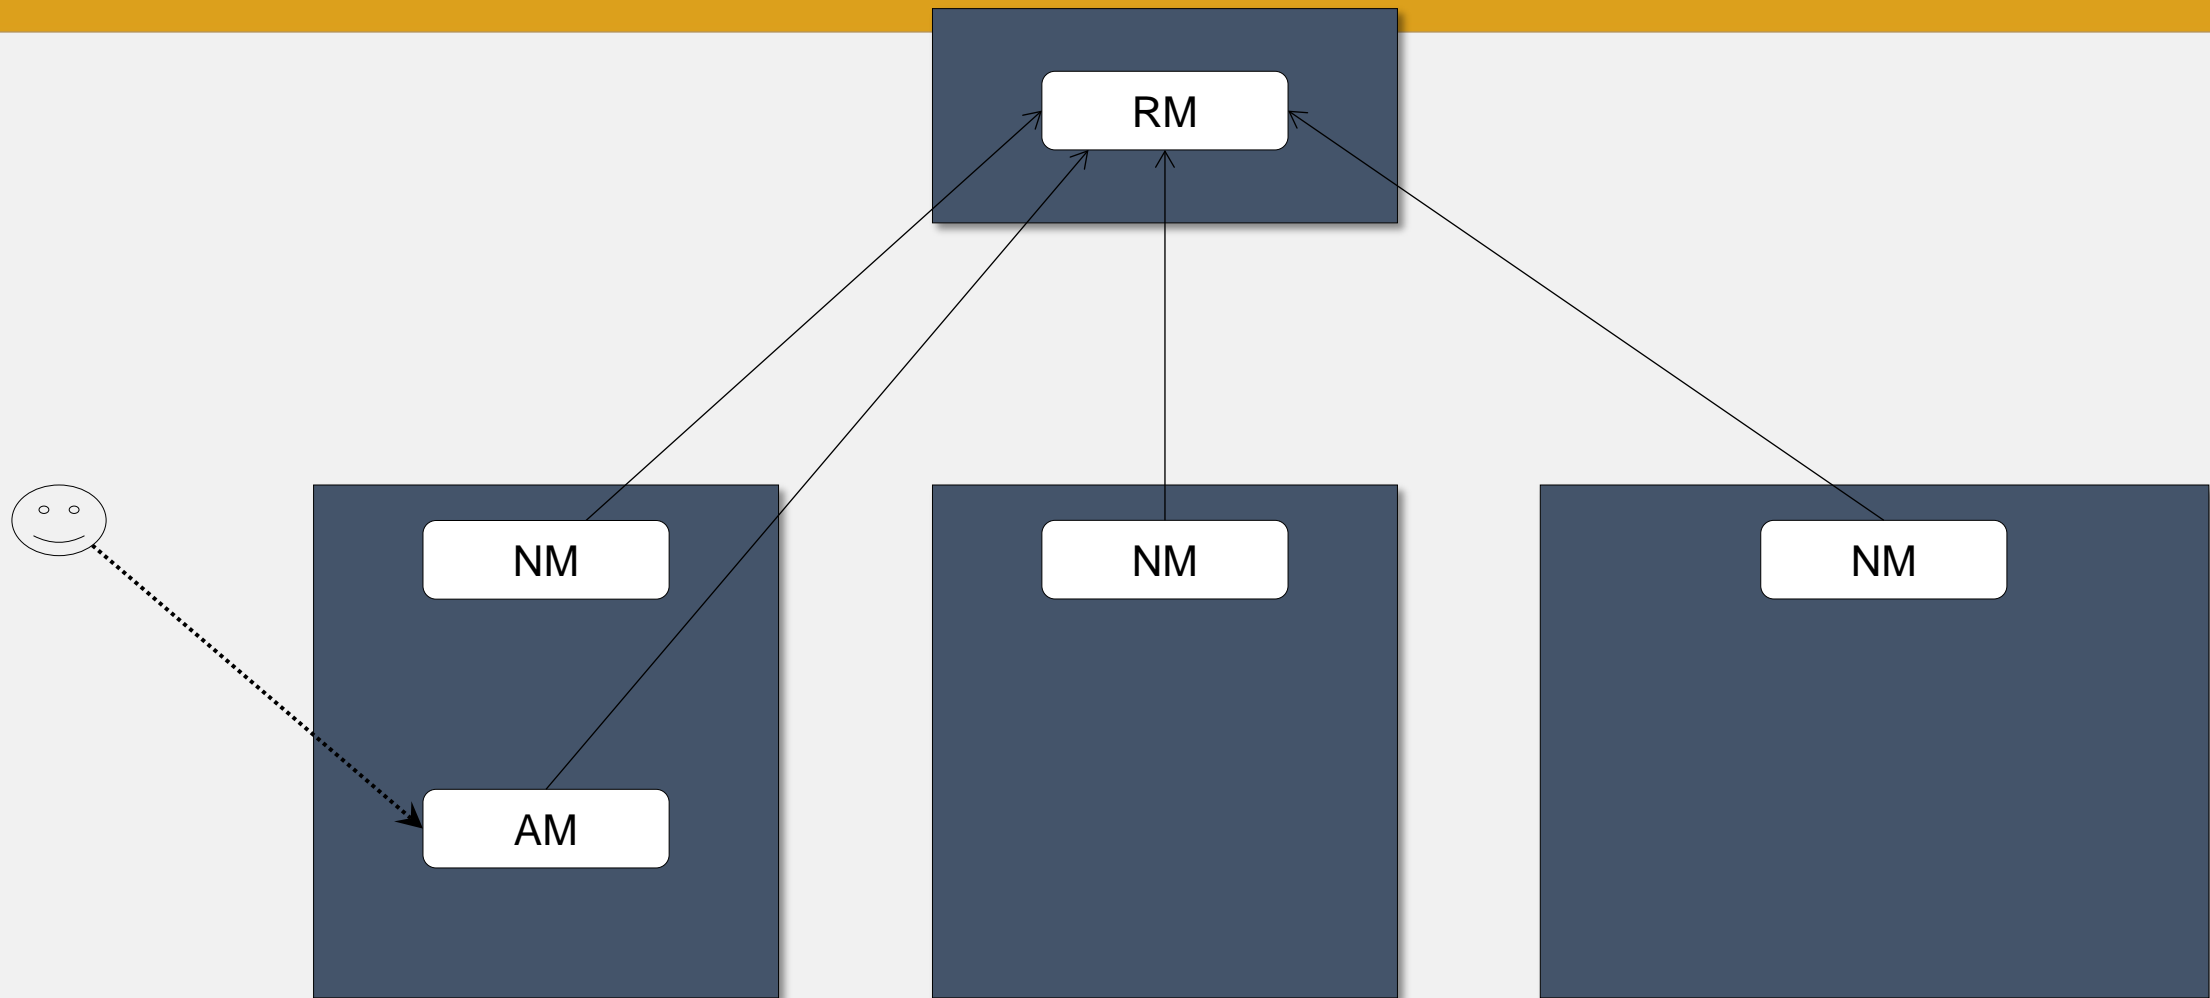
AM informs reducer of shuffle data location for map 1;
reduce task connects to NM and retrieves

MapReduce on YARN



Map 2 completes. AM informs reducer of shuffle data location for map 2 and reducer retrieves

MapReduce on YARN



Reducer completes. AM commits final output, copies job history, cleans staging area, and unregisters from RM before exiting